

# Day-Ahead Power Peak Prediction PoC

## Introduction

This notebook presents a proof-of-concept (PoC) for predicting power peaks using an hourly energy dataset merged with temperature data. Our approach is inspired by the framework proposed in Xia et al. (2024), which focuses on accurately forecasting peak loads at the individual household level. In our case, we adapt the method for an aggregated application (e.g., for malls) by leveraging hourly data.

## Key Components

### 1. Data Aggregation and Preprocessing

- **Aggregation:** We start by aggregating minute-level energy data to an hourly resolution to match the hourly temperature data.
- **Normalization:** Daily load profiles are created and normalized to facilitate pattern extraction and comparison across different days.

### 2. Time Interval Extraction

- The average daily load profile is analyzed to detect local peaks and valleys.
- These peaks and valleys help segment the day into meaningful time intervals (such as morning, afternoon, and evening) where distinct peak load patterns occur.

### 3. Load Decomposition

- **First Stage:** The total daily load is decomposed into:
  - **Base Load:** The stable, underlying consumption.
  - **Peak Load:** The extra consumption above a defined threshold.
- **Second Stage:** The peak load is further broken down into:
  - **Magnitude:** How high the extra consumption is.
  - **Timing:** When during the day the peak occurs.

### 4. Forecasting and Model Comparison

- **Baseline:** We compare a direct forecasting approach (predicting the full daily load without decomposition) as a baseline.
- **Decomposition-Based Approach:** We then forecast the base load and separately predict the peak load's magnitude and timing.
- **Method Comparison:** Additionally, we compare deterministic models (e.g., Gaussian Process Regression, Random Forest, Gradient Boosting) against their probabilistic counterparts (e.g., quantile regression versions) to evaluate how well they capture both the magnitude and the timing of the peaks.

## Rationale

- **Addressing Volatility:**

Traditional time series models (like ARIMA) often struggle with the volatile nature of individual load curves and the "double penalty" problem—where errors in both the magnitude and timing of peaks lead to compounding inaccuracies. Besides such models has a tendency to elvel out the complexe patterns found in hourly and minute domain, making them not suitable for our context.

- **Decomposition Benefits:**

By decomposing the load into base and peak components and further breaking the peak load into its magnitude and timing, the approach allows us to tackle the unpredictable nature of peak events more effectively.

- **Machine Learning Focus:**

The framework from the paper primarily employs machine learning methods (such as Gaussian Process Regression, Random Forest, and Gradient Boosting) which have shown superior performance in capturing non-linear relationships and uncertainty in peak loads compared to classical statistical models.

This notebook outlines the steps from data loading and preprocessing through to feature engineering, model training, and evaluation—demonstrating why the decomposition-based approach, along with a comparison between deterministic and probabilistic methods, provides a more robust solution for power peak prediction than a direct forecasting approach.

---

## Import Libraries and Load Dataset

In this cell, we import the necessary Python libraries and load the combined energy and temperature dataset from "energy\_and\_temperature\_minute\_data.csv." We then convert the timestamp column to datetime.

```
In [46]: # Cell 1: Import necessary libraries and Load the dataset
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.signal import find_peaks

%matplotlib inline

# Load the combined dataset (energy demand + temperature)
data_file = "energy_and_temperature_minute_data.csv"
df = pd.read_csv(data_file, sep=";", low_memory=False)

# Convert the timestamp column to datetime (assuming the column is named 'timestamp_utc')
df["timestamp_utc"] = pd.to_datetime(df["timestamp_utc"], utc=True, errors="coerce")

# Inspect the first few rows and column names
print("Columns in dataset:")
print(df.columns.tolist())
print("\nFirst few rows:")
```

```
print(df.head())
print(df.info())
```

Columns in dataset:

```
['timestamp_utc', 'active_power_W', 'active_power_neg_W', 'reactive_power_VAr',
'reactive_power_neg_VAr', 'phase1_current_A', 'phase2_current_A', 'phase3_current_A',
'phase1_voltage_V', 'phase2_voltage_V', 'phase3_voltage_V', 'meter_id', 'air_temperature']
```

First few rows:

	timestamp_utc	active_power_W	active_power_neg_W	\
0	2025-02-07 13:33:31.488000+00:00	802	0	
1	2025-02-07 13:33:41.628000+00:00	809	0	
2	2025-02-07 13:33:51.571000+00:00	808	0	
3	2025-02-07 13:34:01.563000+00:00	826	0	
4	2025-02-07 13:34:11.498000+00:00	852	0	

	reactive_power_VAr	reactive_power_neg_VAr	phase1_current_A	\
0	16	0	12640	
1	16	0	12670	
2	15	0	12550	
3	14	0	12920	
4	19	0	13060	

	phase2_current_A	phase3_current_A	phase1_voltage_V	phase2_voltage_V	\
0	10940	11220	2330	2320	
1	10980	11420	2330	2320	
2	11010	11460	2330	2320	
3	11270	11590	2330	2320	
4	11750	12140	2330	2320	

	phase3_voltage_V	meter_id	air_temperature
0	2320	6kPJw9QF	2.45
1	2320	6kPJw9QF	2.45
2	2320	6kPJw9QF	2.45
3	2320	6kPJw9QF	2.45
4	2320	6kPJw9QF	2.45

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 126528 entries, 0 to 126527

Data columns (total 13 columns):

#	Column	Non-Null Count	Dtype
0	timestamp_utc	126400 non-null	datetime64[ns, UTC]
1	active_power_W	126528 non-null	int64
2	active_power_neg_W	126528 non-null	int64
3	reactive_power_VAr	126528 non-null	int64
4	reactive_power_neg_VAr	126528 non-null	int64
5	phase1_current_A	126528 non-null	int64
6	phase2_current_A	126528 non-null	int64
7	phase3_current_A	126528 non-null	int64
8	phase1_voltage_V	126528 non-null	int64
9	phase2_voltage_V	126528 non-null	int64
10	phase3_voltage_V	126528 non-null	int64
11	meter_id	126528 non-null	object
12	air_temperature	126528 non-null	float64

dtypes: datetime64[ns, UTC](1), float64(1), int64(10), object(1)

memory usage: 12.5+ MB

None

## Filter Complete Days and Aggregate to Hourly Data

In this cell, we first extract the date from our minute-level timestamps and count the number of records per sensor per day. A full day is defined as having 1440 records (i.e., one reading per minute for 24 hours). We then filter the dataset to keep only those days with complete data. Finally, we aggregate the filtered data to an hourly resolution and report how many days were missing per sensor.

```
In [74]: # Ensure the timestamp is in datetime format and extract date
df["date"] = df["timestamp_utc"].dt.date

# Count records per sensor per day
records_per_day = df.groupby(["meter_id", "date"]).size().reset_index(name="count")

# Define 80% of full day
threshold = 1152

# Identify complete days (those with exactly 1440 records)
complete_days = records_per_day[records_per_day["count"] >= threshold]

print("Number of complete days found per sensor:")
print(complete_days.groupby("meter_id")["date"].count())

# Create a set of complete (meter_id, date) pairs
complete_set = set(zip(complete_days["meter_id"], complete_days["date"]))

# Filter the original dataframe to include only rows from complete days
df_complete = df[df.apply(lambda row: (row["meter_id"], row["date"]) in complete_set, axis=1)]
print("\nOriginal dataset shape:", df.shape)
print("Filtered dataset shape (complete days only):", df_complete.shape)

# Count missing days per sensor:
original_days = df.groupby("meter_id")["date"].nunique().reset_index(name="original_days")
filtered_days = df_complete.groupby("meter_id")["date"].nunique().reset_index(name="filtered_days")
missing_days = pd.merge(original_days, filtered_days, on="meter_id", how="left")
missing_days["missing_days"] = missing_days["original_days"] - missing_days["filtered_days"]
print("\nMissing days per sensor:")
print(missing_days)

# Aggregate the complete data to hourly data
df_complete["timestamp_hour"] = df_complete["timestamp_utc"].dt.floor("H")
hourly_df = df_complete.groupby(["timestamp_hour", "meter_id"]).agg({
    "active_power_W": "mean",
    "active_power_neg_W": "mean",
    "reactive_power_VAr": "mean",
    "reactive_power_neg_VAr": "mean",
    "phase1_current_A": "mean",
    "phase2_current_A": "mean",
    "phase3_current_A": "mean",
    "phase1_voltage_V": "mean",
    "phase2_voltage_V": "mean",
    "phase3_voltage_V": "mean",
    "air_temperature": "mean"
}).reset_index()
```

```
# Rename the aggregated timestamp back to timestamp_utc if desired
hourly_df.rename(columns={"timestamp_hour": "timestamp_utc"}, inplace=True)
print("\nHourly aggregated data:")
print(hourly_df.head())
```

Number of complete days found per sensor:

```
meter_id
6kPJw9QF    35
Jfmwhk2e    18
KGdRbnJc    35
Name: date, dtype: int64
```

Original dataset shape: (126528, 15)

Filtered dataset shape (complete days only): (121427, 15)

Missing days per sensor:

	meter_id	original_days	complete_days	missing_days
0	6kPJw9QF	37	35	2
1	Jfmwhk2e	22	18	4
2	KGdRbnJc	37	35	2

Hourly aggregated data:

	timestamp_utc	meter_id	active_power_W	active_power_neg_W	\
0	2025-02-08 00:00:00+00:00	6kPJw9QF	238.186441	0.0	
1	2025-02-08 00:00:00+00:00	Jfmwhk2e	450.736842	0.0	
2	2025-02-08 00:00:00+00:00	KGdRbnJc	44.771930	0.0	
3	2025-02-08 01:00:00+00:00	6kPJw9QF	213.315789	0.0	
4	2025-02-08 01:00:00+00:00	Jfmwhk2e	463.500000	0.0	

	reactive_power_VAr	reactive_power_neg_VAr	phase1_current_A	\
0	29.406780	0.000000	3799.491525	
1	0.000000	53.421053	5844.385965	
2	40.192982	0.000000	770.175439	
3	29.140351	0.000000	3427.543860	
4	0.000000	51.482759	6151.379310	

	phase2_current_A	phase3_current_A	phase1_voltage_V	phase2_voltage_V	\
0	3117.627119	3463.050847	2335.084746	2326.610169	
1	7095.263158	6070.175439	2396.842105	2398.070175	
2	944.736842	789.298246	2396.842105	2398.070175	
3	2790.175439	3188.245614	2318.596491	2311.403509	
4	7386.379310	5947.068966	2401.724138	2403.103448	

	phase3_voltage_V	air_temperature
0	2326.440678	-0.683333
1	2403.684211	-0.683333
2	2402.456140	-0.683333
3	2311.228070	-0.650000
4	2410.689655	-0.650000

```

C:\Users\saidh\AppData\Local\Temp\ipykernel_43688\187453696.py:33: FutureWarning:
'H' is deprecated and will be removed in a future version, please use 'h' instead.
    df_complete["timestamp_hour"] = df_complete["timestamp_utc"].dt.floor("H")
C:\Users\saidh\AppData\Local\Temp\ipykernel_43688\187453696.py:33: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    df_complete["timestamp_hour"] = df_complete["timestamp_utc"].dt.floor("H")

```

## Analyze and Visualize Data Frequency for All Sensors

In this cell, we loop over all unique sensors (`meter_id`) and compute the time differences between consecutive records for each sensor. We then plot histograms of these time differences to verify the data frequency for each sensor.

```

In [75]: # Get the list of unique sensor IDs
sensor_ids = hourly_df["meter_id"].unique()
n_sensors = len(sensor_ids)

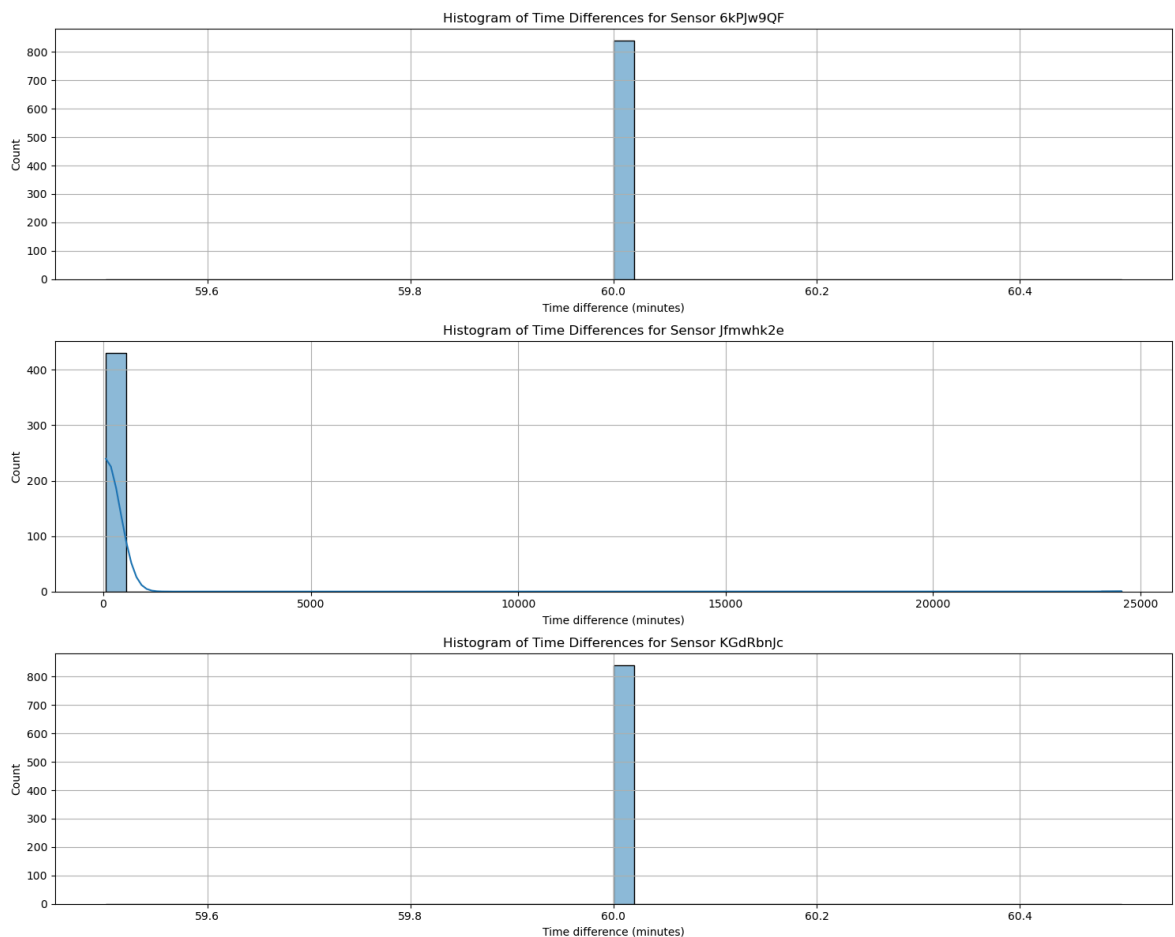
plt.figure(figsize=(15, 4 * n_sensors))

for i, sensor in enumerate(sensor_ids, 1):
    df_filtered_sensor = hourly_df[hourly_df["meter_id"] == sensor].sort_values(
        # Compute time differences in minutes between consecutive records
        df_filtered_sensor["time_diff_min"] = df_filtered_sensor["timestamp_utc"].di

    # Create a subplot for each sensor
    plt.subplot(n_sensors, 1, i)
    sns.histplot(df_filtered_sensor["time_diff_min"].dropna(), bins=50, kde=True)
    plt.xlabel("Time difference (minutes)")
    plt.title(f"Histogram of Time Differences for Sensor {sensor}")
    plt.grid(True)

plt.tight_layout()
plt.show()

```



## Create Daily Load Profiles

We pivot the hourly data to create daily load profiles for each sensor. Each profile will contain 24 hourly values.

```
In [76]: hourly_df["date"] = hourly_df["timestamp_utc"].dt.date

# Pivot to get 24 hourly values per day per sensor (for active power)
daily_profiles = hourly_df.pivot_table(
    index=["date", "meter_id"],
    columns=hourly_df["timestamp_utc"].dt.hour,
    values="active_power_W"
)

# Rename the hourly columns to a consistent format (e.g., hour_0, hour_1, ..., hour_23)
daily_profiles.columns = [f"hour_{col}" for col in daily_profiles.columns]
daily_profiles.reset_index(inplace=True)

print("Sample daily load profiles:")
print(daily_profiles.head())
```

Sample daily load profiles:

	date	meter_id	hour_0	hour_1	hour_2	hour_3	\
0	2025-02-08	6kPJw9QF	238.186441	213.315789	204.724138	205.875000	
1	2025-02-08	Jfmwhk2e	450.736842	463.500000	450.660714	450.775862	
2	2025-02-08	KGdRbnJc	44.771930	44.879310	45.120690	45.228070	
3	2025-02-09	6kPJw9QF	229.206897	211.724138	207.333333	212.913793	
4	2025-02-09	Jfmwhk2e	460.964912	449.578947	458.210526	479.839286	
	hour_4	hour_5	hour_6	hour_7	...	hour_14	\
0	207.333333	242.288136	262.551724	318.964912	...	857.392857	
1	464.339286	527.339286	560.666667	593.859649	...	1172.696429	
2	43.928571	43.105263	43.578947	43.672414	...	96.732143	
3	209.741379	206.535714	215.508772	215.345455	...	279.844828	
4	469.750000	492.245614	484.456140	473.678571	...	477.965517	
	hour_15	hour_16	hour_17	hour_18	hour_19	hour_20	\
0	882.327273	853.689655	745.385965	420.071429	350.327586	313.303571	
1	1146.696429	1094.596491	1085.807018	749.859649	562.875000	482.310345	
2	96.333333	98.189655	97.789474	45.474576	44.931034	45.070175	
3	289.035714	310.157895	304.894737	290.000000	288.666667	291.101695	
4	471.596491	496.473684	544.875000	517.964286	514.689655	502.571429	
	hour_21	hour_22	hour_23				
0	274.500000	272.322034	242.350877				
1	486.779661	491.879310	474.103448				
2	44.551724	43.928571	44.120690				
3	297.333333	254.614035	235.631579				
4	496.913793	501.807018	488.842105				

[5 rows x 26 columns]

## Normalize Daily Load Profiles and Display Average Profiles per Sensor

In this cell, we normalize each day's load profile by dividing the 24 hourly values by the day's maximum load. Then, we compute the average normalized daily load profile for each sensor separately and plot them on the same graph for comparison.

```
In [ ]: hourly_cols = [col for col in daily_profiles.columns if col.startswith("hour_")]

def normalize_profile(row):
    vals = row[hourly_cols].values.astype(float)
    max_val = np.max(vals)
    if max_val > 0:
        return vals / max_val
    else:
        return vals

# Apply normalization to each day's load profile
daily_profiles_normalized = daily_profiles.copy()
daily_profiles_normalized[hourly_cols] = daily_profiles_normalized[hourly_cols].

# Get the unique sensor IDs
sensors = daily_profiles_normalized["meter_id"].unique()

# Plot the average normalized daily load profile for each sensor
plt.figure(figsize=(10, 6))
```




```

for sensor in sensors:
    sensor_data = daily_profiles_normalized[daily_profiles_normalized["meter_id"
    avg_profile_sensor = sensor_data[hourly_cols].mean()
    plt.plot(range(24), avg_profile_sensor, marker="o", label=f"Sensor {sensor}")

plt.title("Average Normalized Daily Load Profile per Sensor")
plt.xlabel("Hour of Day")
plt.ylabel("Normalized Load")
plt.xticks(range(24))
plt.legend()
plt.grid(True)
plt.show()

```



```

In [81]: # Cell X+2: Detect peaks and valleys for each sensor's average normalized daily
from scipy.signal import find_peaks

# Get unique sensor IDs from your normalized daily profiles
sensors = daily_profiles_normalized["meter_id"].unique()
n_sensors = len(sensors)

# Create subplots to visualize each sensor separately
fig, axes = plt.subplots(n_sensors, 1, figsize=(10, 4 * n_sensors), sharex=True)

# Loop over each sensor to compute and plot its average profile with detected peaks
for i, sensor in enumerate(sensors):
    sensor_data = daily_profiles_normalized[daily_profiles_normalized["meter_id"
    # Compute the average normalized daily load profile for the sensor
    avg_profile_sensor = sensor_data[hourly_cols].mean()

    # Detect peaks in the average profile
    peaks, _ = find_peaks(avg_profile_sensor, distance=1)
    # Detect valleys (by finding peaks in the inverted profile)
    valleys, _ = find_peaks(-avg_profile_sensor, distance=1)

    # Plot on the appropriate subplot
    ax = axes[i] if n_sensors > 1 else axes
    ax.plot(range(24), avg_profile_sensor, marker="o", label="Avg Load")

```

```

ax.plot(peaks, avg_profile_sensor[peaks], "x", markersize=10, label="Peaks")
ax.plot(valleys, avg_profile_sensor[valleys], "v", markersize=10, label="Valleys")
ax.set_title(f"Sensor {sensor}: Average Normalized Daily Load Profile")
ax.set_xlabel("Hour of Day")
ax.set_ylabel("Normalized Load")
ax.set_xticks(range(24))
ax.legend()
ax.grid(True)

plt.tight_layout()
plt.show()

```

C:\Users\saidh\AppData\Local\Temp\ipykernel\_43688\2760446431.py:25: FutureWarning: Series.\_\_getitem\_\_ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`

```
ax.plot(peaks, avg_profile_sensor[peaks], "x", markersize=10, label="Peaks")
```

C:\Users\saidh\AppData\Local\Temp\ipykernel\_43688\2760446431.py:26: FutureWarning: Series.\_\_getitem\_\_ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`

```
ax.plot(valleys, avg_profile_sensor[valleys], "v", markersize=10, label="Valleys")
```

C:\Users\saidh\AppData\Local\Temp\ipykernel\_43688\2760446431.py:25: FutureWarning: Series.\_\_getitem\_\_ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`

```
ax.plot(peaks, avg_profile_sensor[peaks], "x", markersize=10, label="Peaks")
```

C:\Users\saidh\AppData\Local\Temp\ipykernel\_43688\2760446431.py:26: FutureWarning: Series.\_\_getitem\_\_ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`

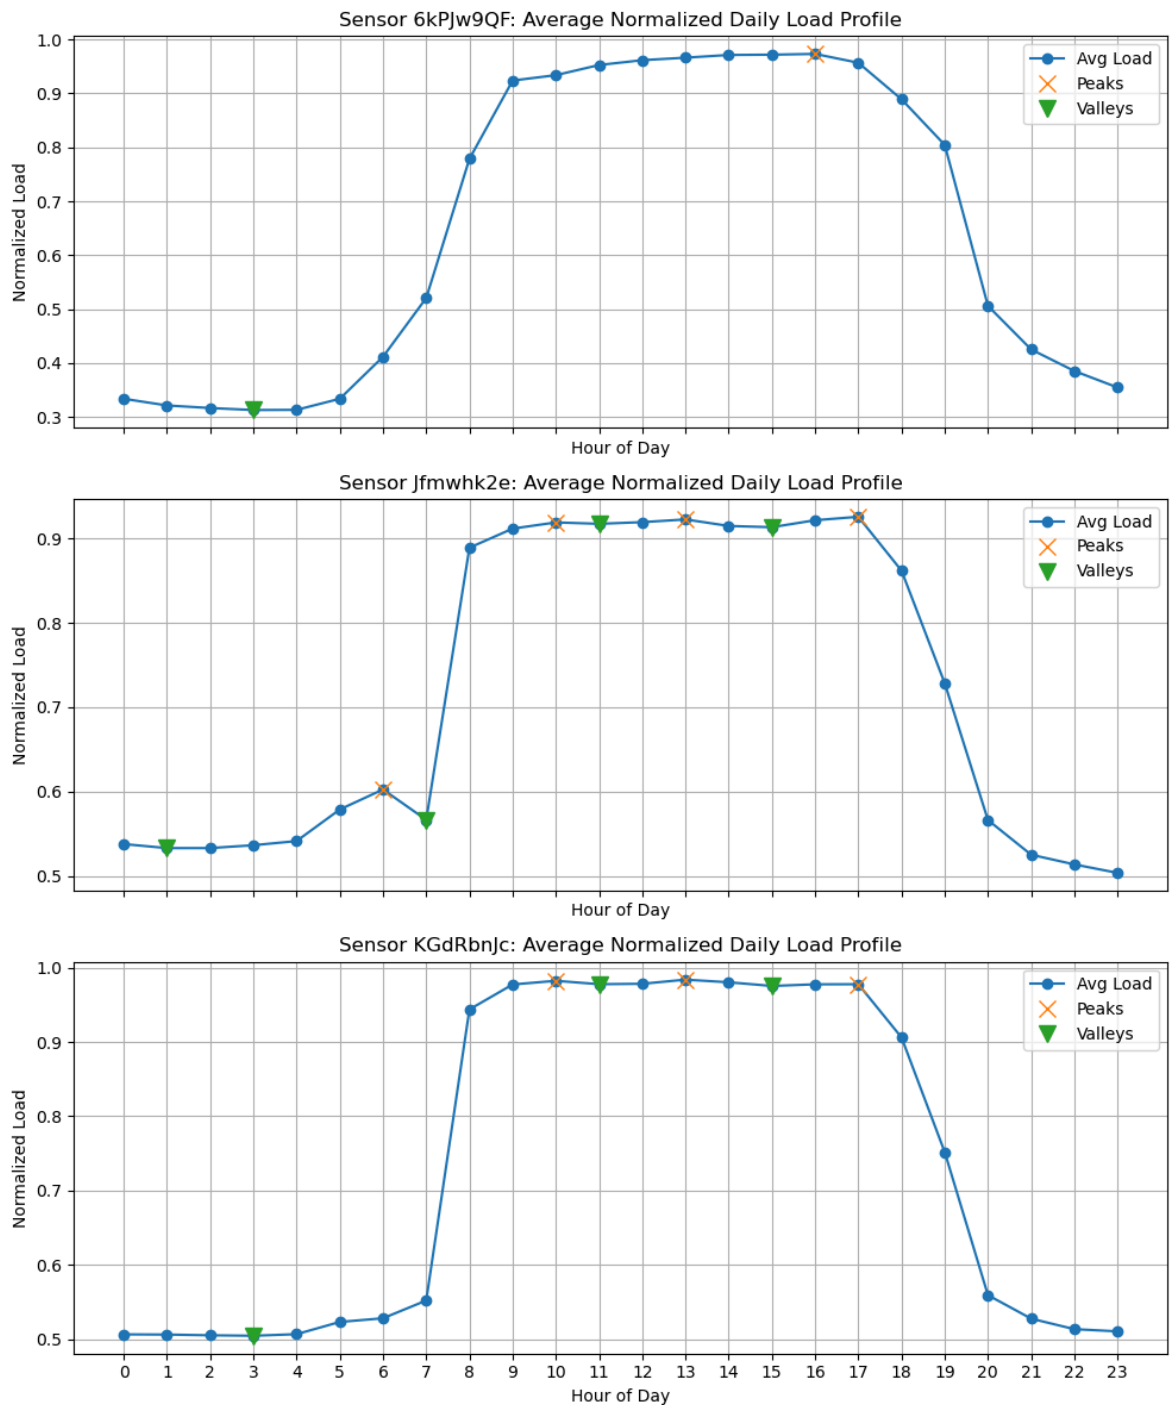
```
ax.plot(valleys, avg_profile_sensor[valleys], "v", markersize=10, label="Valleys")
```

C:\Users\saidh\AppData\Local\Temp\ipykernel\_43688\2760446431.py:25: FutureWarning: Series.\_\_getitem\_\_ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`

```
ax.plot(peaks, avg_profile_sensor[peaks], "x", markersize=10, label="Peaks")
```

C:\Users\saidh\AppData\Local\Temp\ipykernel\_43688\2760446431.py:26: FutureWarning: Series.\_\_getitem\_\_ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`

```
ax.plot(valleys, avg_profile_sensor[valleys], "v", markersize=10, label="Valleys")
```



## Decompose Daily Load into Base Load and Peak Load

In this cell, we decompose each sensor's normalized daily load profile into two parts:

- **Base Load:** The portion of the load up to a boundary (here, using the median).
- **Peak Load:** The excess load above that boundary.

We also extract the hour at which the peak (excess) is highest. This mirrors the paper's first-stage decomposition.

```
In [82]: def decompose_load(row):
# Extract the 24 hourly values from the normalized profile
values = row[hourly_cols].values.astype(float)
```

```

# Define the base-peak boundary as the median of the daily profile
boundary = np.median(values)
# Base Load: every hour's consumption capped at the boundary
base_load = np.minimum(values, boundary)
# Peak Load: the extra consumption above the boundary
peak_load = np.maximum(values - boundary, 0)
# Return the total base load, total peak load, and the hour when peak load is
return pd.Series({
    "total_base_load": np.sum(base_load),
    "total_peak_load": np.sum(peak_load),
    "peak_load_hour": np.argmax(peak_load) # returns an integer 0-23
})

# Apply decomposition to each daily profile (using the normalized profiles)
decomposition = daily_profiles_normalized.apply(decompose_load, axis=1)
# Merge the decomposition results with the original daily profiles
daily_profiles = pd.concat([daily_profiles, decomposition], axis=1)

print("Sample of daily profiles with decomposition:")
print(daily_profiles.head())

```

Sample of daily profiles with decomposition:

	date	meter_id	hour_0	hour_1	hour_2	hour_3	\
0	2025-02-08	6kPJw9QF	238.186441	213.315789	204.724138	205.875000	
1	2025-02-08	Jfmwhk2e	450.736842	463.500000	450.660714	450.775862	
2	2025-02-08	KGdRbnJc	44.771930	44.879310	45.120690	45.228070	
3	2025-02-09	6kPJw9QF	229.206897	211.724138	207.333333	212.913793	
4	2025-02-09	Jfmwhk2e	460.964912	449.578947	458.210526	479.839286	

	hour_4	hour_5	hour_6	hour_7	...	hour_17	\
0	207.333333	242.288136	262.551724	318.964912	...	745.385965	
1	464.339286	527.339286	560.666667	593.859649	...	1085.807018	
2	43.928571	43.105263	43.578947	43.672414	...	97.789474	
3	209.741379	206.535714	215.508772	215.345455	...	304.894737	
4	469.750000	492.245614	484.456140	473.678571	...	544.875000	

	hour_18	hour_19	hour_20	hour_21	hour_22	hour_23	\
0	420.071429	350.327586	313.303571	274.500000	272.322034	242.350877	
1	749.859649	562.875000	482.310345	486.779661	491.879310	474.103448	
2	45.474576	44.931034	45.070175	44.551724	43.928571	44.120690	
3	290.000000	288.666667	291.101695	297.333333	254.614035	235.631579	
4	517.964286	514.689655	502.571429	496.913793	501.807018	488.842105	

	total_base_load	total_peak_load	peak_load_hour
0	7.946565	5.406949	15.0
1	10.554154	4.597310	13.0
2	10.935489	5.222858	16.0
3	18.389216	1.310009	16.0
4	20.679490	0.504428	17.0

[5 rows x 29 columns]

## Feature Engineering for Peak Load Forecasting

Next, we construct additional features that will help predict peak load characteristics. In this cell we:

- Create lag features (e.g., the previous day's, two days ago, and one week ago's total peak load) for each sensor.
- Merge daily average temperature information from the hourly data.

```
In [ ]: # Ensure the daily_profiles is sorted by sensor and date
daily_profiles.sort_values(by=["meter_id", "date"], inplace=True)
daily_profiles["date"] = pd.to_datetime(daily_profiles["date"])

# Create lag features for the total peak load
daily_profiles["lag1_peak"] = daily_profiles.groupby("meter_id")["total_peak_load"].shift(1)
daily_profiles["lag2_peak"] = daily_profiles.groupby("meter_id")["total_peak_load"].shift(2)
daily_profiles["lag7_peak"] = daily_profiles.groupby("meter_id")["total_peak_load"].shift(7)

# Build a daily temperature feature from the hourly data: compute daily average
hourly_df["date"] = hourly_df["timestamp_utc"].dt.date
daily_temp = hourly_df.groupby("date")["air_temperature"].mean().reset_index()
daily_temp["date"] = pd.to_datetime(daily_temp["date"])

# Merge the daily temperature with the daily profiles based on date
daily_profiles = pd.merge(daily_profiles, daily_temp, on="date", how="left")

print("Daily profiles with additional features:")
print(daily_profiles.head())
```

Daily profiles with additional features:

	date	meter_id	hour_0	hour_1	hour_2	hour_3	\
0	2025-02-08	6kPJw9QF	238.186441	213.315789	204.724138	205.875000	
1	2025-02-09	6kPJw9QF	229.206897	211.724138	207.333333	212.913793	
2	2025-02-10	6kPJw9QF	224.543860	209.271186	206.842105	202.000000	
3	2025-02-11	6kPJw9QF	243.350877	229.192982	226.033898	229.824561	
4	2025-02-12	6kPJw9QF	230.232143	221.224138	217.438596	208.087719	

	hour_4	hour_5	hour_6	hour_7	...	hour_21	\
0	207.333333	242.288136	262.551724	318.964912	...	274.500000	
1	209.741379	206.535714	215.508772	215.345455	...	297.333333	
2	209.368421	250.830508	304.052632	438.694915	...	294.771930	
3	237.500000	244.241379	328.288136	453.298246	...	272.491228	
4	213.322034	224.661017	313.385965	432.913793	...	277.338983	

	hour_22	hour_23	total_base_load	total_peak_load	peak_load_hour	\
0	272.322034	242.350877	7.946565	5.406949	15.0	
1	254.614035	235.631579	18.389216	1.310009	16.0	
2	282.241379	260.385965	11.735416	3.357954	15.0	
3	249.465517	241.543860	12.589608	3.053788	10.0	
4	248.724138	235.327273	11.777175	3.539961	16.0	

	lag1_peak	lag2_peak	lag7_peak	air_temperature
0	NaN	NaN	NaN	-0.851736
1	5.406949	NaN	NaN	-5.063819
2	1.310009	5.406949	NaN	-3.871528
3	3.357954	1.310009	NaN	-2.950694
4	3.053788	3.357954	NaN	-4.679306

[5 rows x 33 columns]

## Model Training for Peak Load Magnitude Prediction

Here, we train a simple regression model to predict the total peak load (magnitude) using our lag features and daily average temperature. This model represents the deterministic forecasting part for the peak magnitude.

```
In [84]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error

# Drop rows with missing values in lag features or target variable
model_data = daily_profiles.dropna(subset=["lag1_peak", "lag2_peak", "lag7_peak"])

# Define features and target
features = ["lag1_peak", "lag2_peak", "lag7_peak", "air_temperature"]
target = "total_peak_load"

X = model_data[features]
y = model_data[target]

# Split into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_

# Train a RandomForestRegressor
model_magnitude = RandomForestRegressor(n_estimators=100, random_state=42)
model_magnitude.fit(X_train, y_train)

# Make predictions and evaluate
preds_magnitude = model_magnitude.predict(X_test)
mae_magnitude = mean_absolute_error(y_test, preds_magnitude)
print(f"Mean Absolute Error (Peak Load Magnitude): {mae_magnitude:.2f}")
```

Mean Absolute Error (Peak Load Magnitude): 0.63

## Model Training for Peak Load Timing Prediction

Now, we set up another model to predict the timing of the peak load (i.e., the hour when the peak occurs). In this example, we treat it as a regression problem (predicting an hour between 0 and 23). You might also consider a classification approach.

```
In [86]: # Cell X+6: Train a model to predict peak_load_hour (timing)
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error

# Drop rows with missing values for the necessary features and target
model_data_time = daily_profiles.dropna(subset=["lag1_peak", "lag2_peak", "lag7_

X_time = model_data_time[features] # Using the same features as before
y_time = model_data_time["peak_load_hour"]

X_train_t, X_test_t, y_train_t, y_test_t = train_test_split(X_time, y_time, test

# Train a RandomForestRegressor for timing
model_timing = RandomForestRegressor(n_estimators=100, random_state=42)
model_timing.fit(X_train_t, y_train_t)

# Predict and evaluate timing error (MAE in hours)
preds_timing = model_timing.predict(X_test_t)
```

```
mae_timing = mean_absolute_error(y_test_t, preds_timing)
print(f"Mean Absolute Error (Peak Load Timing in hours): {mae_timing:.2f}")
```

Mean Absolute Error (Peak Load Timing in hours): 2.83

In [ ]:

## Evaluation and Visualization

Finally, we visualize the performance of our models. Here we plot:

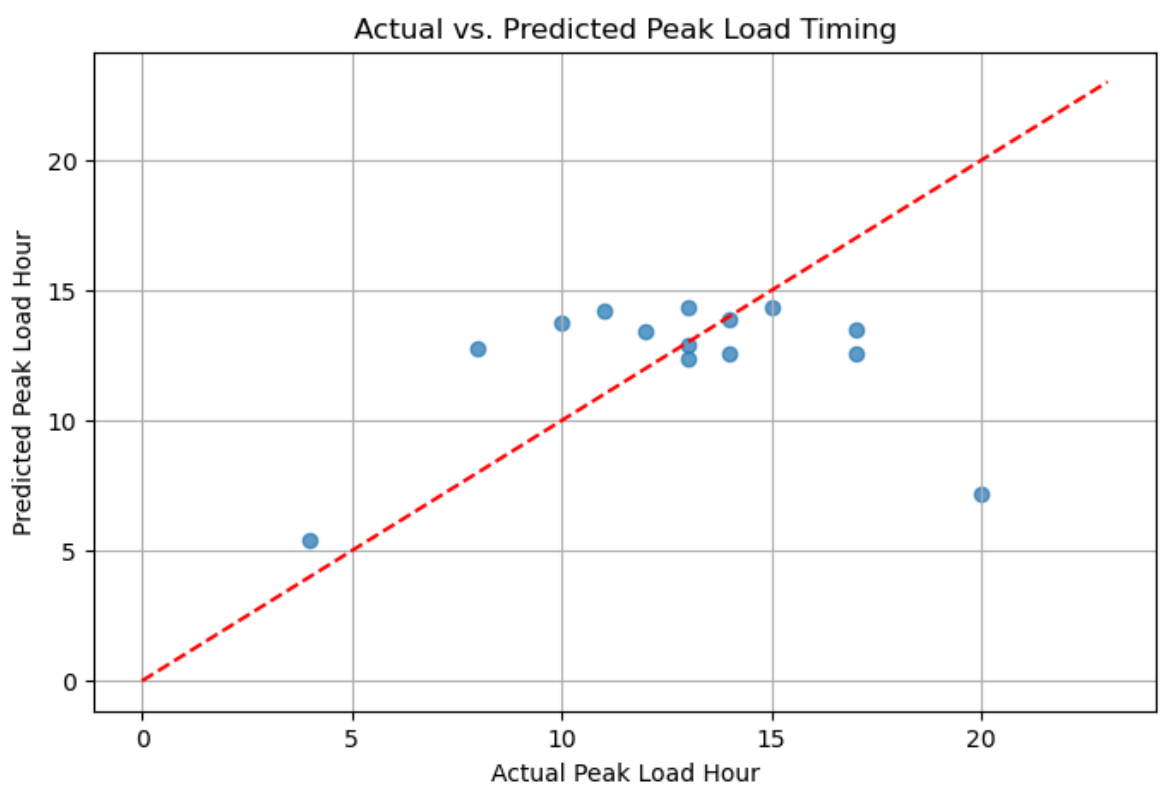
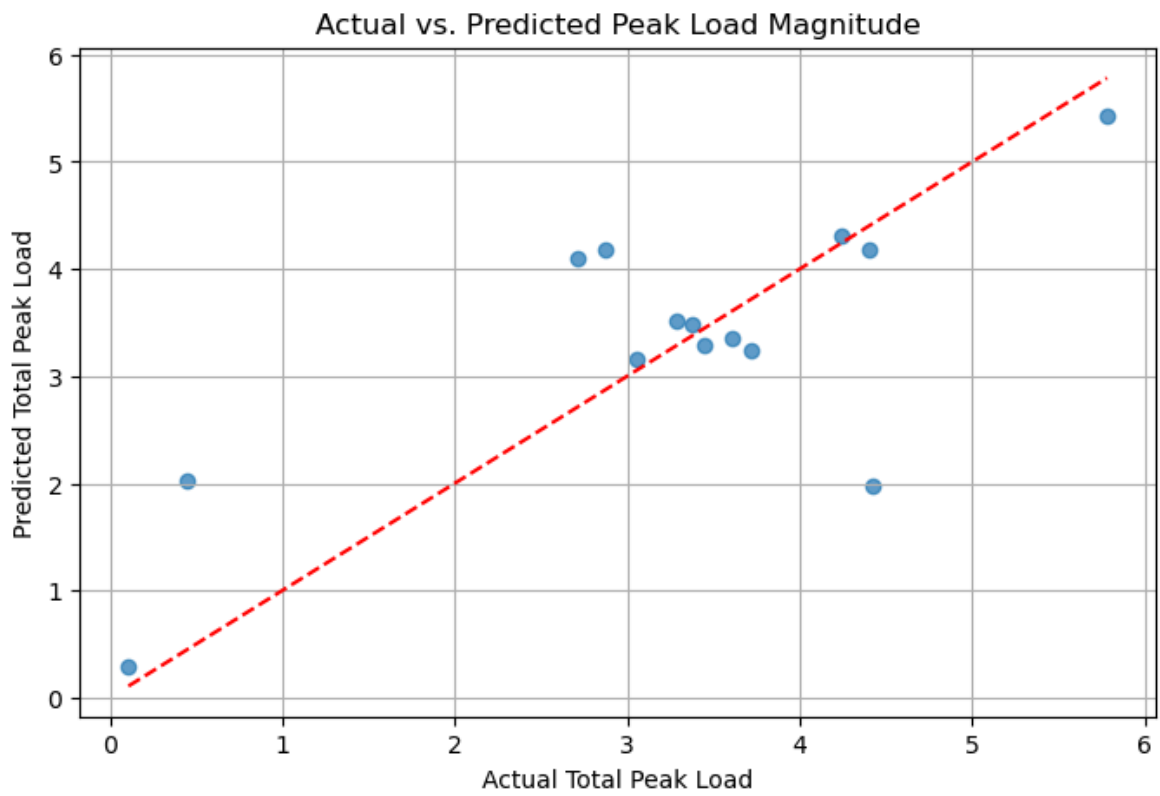
- Actual vs. predicted peak load magnitude.
- Actual vs. predicted peak load timing.

This gives us an initial assessment of how well our models capture the peak characteristics.

```
In [87]: # Plot for Peak Load Magnitude
plt.figure(figsize=(8, 5))
plt.scatter(y_test, preds_magnitude, alpha=0.7)
plt.xlabel("Actual Total Peak Load")
plt.ylabel("Predicted Total Peak Load")
plt.title("Actual vs. Predicted Peak Load Magnitude")
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], "r--")
plt.grid(True)
plt.show()

# Plot for Peak Load Timing
plt.figure(figsize=(8, 5))
plt.scatter(y_test_t, preds_timing, alpha=0.7)
plt.xlabel("Actual Peak Load Hour")
plt.ylabel("Predicted Peak Load Hour")
plt.title("Actual vs. Predicted Peak Load Timing")
plt.plot([0, 23], [0, 23], "r--")
plt.grid(True)
plt.show()
```

<Figure size 800x500 with 0 Axes>



In [ ]: